

I/O Systems

Operating Systems – EDA093/DIT401

Vincenzo Gulisano
vincenzo.gulisano@chalmers.se



CHALMERS
UNIVERSITY OF TECHNOLOGY



UNIVERSITY OF
GOTHENBURG

Reading instructions

- Chapter 5.1 to 5.3

Agenda

- Introduction
- I/O Hardware
- Application I/O Interface
- Kernel I/O Subsystem
- Performance

Agenda

- Introduction
- I/O Hardware
- Application I/O Interface
- Kernel I/O Subsystem
- Performance


Introduction

- Main jobs of a computer: I/O & processing.
- **Role of OS:** manage & control I/O operations & I/O devices

OBJECTIVES:

- Explore the structure of an OS's I/O subsystem
- Describe the principles and complexities of I/O hardware
- Explain performance aspects of I/O hardware and software

Introduction



I/O challenge: control the devices in a given architecture as effectively as possible

Why is it a challenge?

1. **Variability** of I/O devices (in their function & speed)
Varied methods are needed to control them.
2. Software and hardware interfaces are becoming more standardized, while the variety of devices continues to increase
3. Different details & oddities of devices

How does the OS address the challenge?

- dedicated I/O subsystem in the kernel
- **device drivers** that provide a **device access** interface to the I/O subsystem

Agenda

- Introduction
- I/O Hardware
- Application I/O Interface
- Kernel I/O Subsystem
- Performance


I/O Hardware

We have different types of hardware usually available in a computer:

- Storage devices (disks)
- Transmission devices (network cards)
- Human-interface devices (screen, keyboard, mouse)

All of them exchange information by sending signals (via cable or air)

How do they communicate with the CPU, then?



I/O Hardware

Common concepts:

- **Port:** the connection point (e.g., serial port)
- **Bus:** a set of wires & a protocol that specifies a set of messages that can be sent on the wires.
 - Messages conveyed by patterns of electrical voltages with different timings.
 - Buses vary in signaling methods, speed, throughput & connection methods.
- When device A connects to device B and device B connects to device C...
We say A,B,C,... form a **daisy chain**
 - In some systems, daisy-chained devices share a bus or pass bus signals from one device to the next

Typical PC bus structure (legacy)

Important concepts:

- **PCI bus:** connects the processor-memory subsystem to the fast devices.
- **Expansion bus:** connects relatively **slow devices**.
- **Controller:** a collection of electronics that can operate a port, a bus or a device.


Example 1: Serial-port controller

- A single chip (on the host) that controls the signals on the wires of a serial port.

Example 2: SCSI controller

- It may contain a processor, microcode, and private memory for processing SCSI protocol messages. It may be implemented on an adapter card or integrated into a device or motherboard.

Example of a legacy PC bus structure



I/O Hardware



How can the processor give commands & data to a controller to accomplish an I/O transfer?

1st Way of Communication


- Processor reads & writes bit patterns in the controllers' registers.
 - Use of special I/O instructions that specify the transfer of a byte or word to an I/O port address.
 - I/O instruction triggers bus lines to select the proper device & move bits into or out of a device register.

2nd Way of Communication: Memory-mapped I/O


- Device-control registers are mapped into the address space of the processor.
- CPU executes I/O requests using the standard data transfer instructions to read and write the device-control registers.

→ Some systems use both ways of communication.

Processor reads & writes bit patterns in the controllers' registers




Memory-mapped I/O



When I want to write a byte to a register of a certain controller I use an ordinary load or store instruction. Address-decoding hardware routes the access to RAM or to a device according to the addressed region.

Example of Memory-mapped I/O: graphics controller



CPU



motherboard



Graphics controller

Graphics controller: has large **memory-mapped** region to hold screen contents.
Process **sends output** to the screen by writing data into the memory-mapped region.
Controller generates **screen image** based on the contents of the memory.

Device I/O Port Locations on PCs

Example legacy x86 I/O-port address ranges

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)

A few fixed I/O-port ranges remain for compatibility, but modern devices generally use dynamically assigned memory-mapped I/O regions.

I/O Hardware

I/O port typically consists of:

- Data-in register is read by the host to get input (e.g. 1 to 4 bytes in size)
- Data-out register is written by the host to send output
- Status register: contains bits that can be read by the host. Indicates states such as:
 - whether the current command has completed
 - whether a byte is available to be read
 - whether a device error has occurred
- Control register: to start a command or to change the mode of a device.

Example 1: a certain bit in the control register of a serial port chooses between full-duplex & half-duplex communication.

Example 2: another bit sets the word length to 7 or 8 bits. Other bits select one of the speeds supported by the serial port

Interaction between host and controller

- Complete interaction protocol between the host & a controller can be intricate.
 - **busy-waiting / polling:** the host repeatedly reads the controller's status register until the required condition holds
 - **interrupt:** the controller notifies the CPU of an event such as input availability, operation completion, or an error

Polling – Handshaking procedure


- 2 bits are used to coordinate the producer-consumer relationship between the controller & the host.
- Controller indicates its **state** through the **busy** bit in the **status** register.
- Host **signals** its wishes via the command-ready bit in the command register.
Command ready bit: if it is set, a command is available for the controller
- **busy-waiting / polling:** it is a loop, reading the status register over and over until the busy bit becomes clear.
Reasonable when the device and its controller respond quickly
Inefficient when the CPU repeatedly polls while the device remains busy
- **In such cases:** hardware controller should notify the CPU when the device becomes **ready** for service → this hardware mechanism is called **interrupt**.

Interrupts

Basic mechanism (simplified)

1. The interrupt mechanism delivers an interrupt request to the CPU, for example through an interrupt line/controller or a message-signaled interrupt.
2. CPU senses the interrupt request line after executing every instruction.
3. If yes, CPU performs a save state & jumps to the interrupt handler routine at a fixed address in memory.
4. Interrupt handler (IH): What was the cause of the interrupt?
5. The handler performs the necessary processing, restores the saved state, and executes a return-from-interrupt instruction.

Interrupts



Interrupts - Dictionary

- The device controller raises an interrupt
-by asserting a signal on the interrupt request line
- CPU takes/accepts the interrupt
- CPU dispatches it to the interrupt handler
- Handler clears the interrupt by servicing the device

Interrupts

Sophisticated Interrupt Handling

In a modern OS: we need sophisticated interrupt handling features

- Ability to **defer interrupt handling** during critical processing
- Efficient way to locate the proper interrupt handler for a device without asking all devices to see which one raised the interrupt
- Multiple interrupt-priority levels allow **higher-priority** events to be handled before **lower-priority** events
- In modern OS these provided by **CPU & interrupt controller hardware**

Interrupts

Two interrupt request lines (this is actually architecture specific):

- **Nonmaskable interrupt:** reserved for events such as unrecoverable memory errors.
- **Maskable interrupt:** may be temporarily masked by the CPU or interrupt controller while a short critical section is executed.

Interrupt mechanism accepts an **address**: a **number** that selects a **specific interrupt-handling routine** from a small set.

→ This address is an **offset** in a table called interrupt vector

This vector contains the **memory addresses** of **specialized interrupt handlers**.

Interrupts – Intel Pentium processor event-vector table

nonmaskable

vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19–31	(Intel reserved, do not use)
32–255	maskable interrupts

Interrupts

OS & interrupts interaction

- A modern OS interacts with the interrupt mechanism in **several ways**.
- At boot time, **probes** the buses to see which devices are present and **installs/configures the corresponding interrupt handlers**.
- During I/O, device controllers **raise interrupts** when they are ready for service. These interrupts may mean:
 - output has completed
 - input data is available
 - or a failure has been detected

Direct Memory Access

Direct Memory Access

- Devices with large transfers (e.g., disk drive).
Not a good idea to use an expensive **general-purpose** processor to watch status bits & feed data into a controller register → **programmed I/O** process.
- Use a DMA controller instead
- After the CPU/driver configures the operation, the DMA engine transfers the block between the device and memory without CPU instructions moving each byte or word.

DMA transfer


- To initiate a DMA transfer, host writes a DMA command block into memory.
- Block contains:
 - a pointer to the source of the transfer
 - a pointer to the destination of the transfer
 - a count of the number of bytes to be transferred.

Direct Memory Access

Simplified Handshaking between DMA controller & device controller

- Performed via a pair of wires: DMA request & DMA-acknowledge.
- A word of data available for transfer → device controller places a signal on the DMA-request wire
- DMA controller:
 - places the **desired memory address** on the memory-address wires.
 - places a signal on the DMA-acknowledge wire.
- Device controller receives the **DMA-acknowledge** signal & transfers the word of data to memory & removes the **DMA-request** signal.

Six Step Process to Perform DMA Transfer




Agenda

- Introduction
- I/O Hardware
- Application I/O Interface
- Kernel I/O Subsystem
- Performance

Application I/O Interface

- **Goals:**
 - **Abstract away** the detailed differences in I/O devices by identifying a few general kinds.
 - Provide general interfaces that encapsulate common device behaviors
- Device-driver layer **hides differences** among I/O controllers from from the generic kernel I/O subsystem.
- Make I/O subsystem **independent** of the hardware:
 - Simplifies the job of the OS.
 - Benefits the hardware manufacturers
 - Standard interfaces allow a compatible peripheral to be attached without changing the generic I/O subsystem, provided that an appropriate driver already exists
 - **BUT:** each OS has its own standards for the device-driver interface.

Application I/O Interface



Application I/O Interface

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read-write	CD-ROM graphics controller disk

Devices vary in many dimensions:

- **Character-stream** or **block**: bytes one by one or blocks of bytes
- **Sequential** or **random access**: fixed order or random
- **Synchronous** or **asynchronous**: data transfers with predictable response times or not
- **Sharable** or **dedicated**: used concurrently by several processes or not
- **Speed of operation**: speeds range from a few bytes to gigabytes per sec
- **I/O direction**: some devices support both reading and writing; others support only one direction.

Block and Character Devices

Block devices include disk drives

Block-device interface provides the operations needed to access disks and other block-oriented devices

- Commands include `read()`, `write()`, `seek()`
- Memory-mapped file access possible

Character devices

- Include keyboards, mouse, serial ports
- Commands include `get()`, `put()`
- Libraries layered on top allow line editing

Network Devices

Network Devices

- Network I/O differs significantly from disks, the interface is not `read()`, `write()` and `seek()`
- Unix & Windows include `socket` interface
- System calls in the socket interface enable an application to
 - Create a socket
 - Connect to a remote socket

Socket Interface

- Connect local socket to a remote address
- Listen remote application to plug into the local socket
- Send & receive packets over the connection
- `select()` → return which sockets have a packet waiting to be received

Clocks & Timers

- Give current time.
- Give elapsed time.
- Set a timer to **trigger** operation X at time T.
- **Programmable hardware interval timer:** hardware to measure elapsed time & trigger operations
- **Example:** wait a certain amount of time & then generate an interrupt
Cancel operations preceding **too slowly**
- **Usual interrupt rate:** between 18 & 60 ticks per sec.
- Hardware clock constructed from a high-frequency counter.

Blocking & Nonblocking I/O

Blocking

- Process suspended until I/O completed
- Moved from run queue to wait queue
- Later moved back to run queue & resumes
- Easy to use & understand
- Insufficient for some needs

Nonblocking

- Example: a user interface that receives keyboard & mouse input while processing & displaying data on screen.
- Implemented via multi-threading (a thread does the blocking part, another processes results)
- Returns quickly with a return value indicating how many bytes were transferred.

Blocking & Nonblocking I/O

Asynchronous

- Process runs while I/O executes
- Returns immediately without waiting for I/O to complete
- Difficult to use
- I/O subsystem signals process when I/O completed

Agenda

- Introduction
- I/O Hardware
- Application I/O Interface
- Kernel I/O Subsystem
- Performance

Kernel I/O subsystem

- Provides many services related to I/O, including:
 - I/O scheduling
 - Buffering / Caching / Spooling
 - Error handling
 - I/O protection

I/O Scheduling

Goal: schedule I/O requests in a **good order** before executing them

- **Improve** overall performance
 - Share device access **fairly** among processes
 - **Reduce** the average waiting time for I/O to complete
-
- Some OS maintain a wait queue of requests for each device
 - Some OS try **fairness**


Example

- Application 1 needs a block near the end of a disk
- Application 2 needs a block near the beginning
- Application 3 needs a block near the middle of the disk

Best Scheduling: 2,3,1

Device status table

- When a kernel supports **asynchronous** I/O and might schedule requests in different orders, it must be able to keep track of many I/O requests




Buffering

- **Another way to improve performance** → Use storage space in main memory or on disk via buffering, caching, spooling

Buffering

- Buffer: memory area that stores data being transferred between two devices or between a device & an application.
- Example:
 - File received via modem for storage on the hard disk.
 - Buffer created in main memory to accumulate the bytes received from the modem.
 - Buffer full → write data on the disk
 - Disk write **not instantaneous**: two buffers are needed → Double buffering

Buffering – Differences in device speeds



Enormous differences in device speeds for typical hardware!

Sun Enterprise 6000 device-transfer rates (logarithmic)

Caching – Spooling

Caching

- Cache: holds a copy on faster storage of an item that resides elsewhere
- Easier access → Key to performance

Spooling

- Hold output for a device
- If device can serve only one request at a time (e.g. printing)
Devices that cannot accept interleaved data streams.

Error Handling

Error Handling

- Devices and I/O transfers can fail in many ways.
 - **Transient reasons:** a network becomes overloaded
 - **Permanent reasons:** a disk controller becomes defective
- OS can recover from disk read, device unavailable, transient write failures
- Most return an error number or code when I/O request fails
- System error logs hold problem reports

I/O Protection

- **Necessary:** User access may attempt to disrupt normal operation via illegal I/O instructions.
- All I/O instructions defined to be **privileged** → Users cannot issue them directly
- To do I/O a user program **executes** a system call
- OS **checks** if request is **valid**
- If yes it is executed.
- Memory-mapped & I/O port memory locations must be **protected**

Agenda

- Introduction
- I/O Hardware
- Application I/O Interface
- Kernel I/O Subsystem
- Performance


Performance

I/O has major impact in system performance:

Why?

- Places heavy demands on CPU to execute device driver
- Schedule processes fairly & efficiently
- Context switches due to interrupts
- ...


Example of performance: remote login



How to improve performance?

- Reduce number of context switches
- Reduce number of times data must be copied in memory (between device & application)
- Reduce frequency of interrupts
Use large transfers, smart controllers
- Use DMA
- Balance CPU, memory, bus and I/O performance for highest throughput

Performance



Questions

Vincenzo Gulisano

When using memory-mapped I/O...

1. The same address can be shared for RAM and I/O device
2. The same address cannot be shared for RAM and I/O device

The address of an interrupt is used to...

1. Detect if the interrupt is maskable or not
2. Find its interrupt-handling routine

2 page faults interrupts can be served in parallel by the same interrupt-handling routine

1. Yes
2. No

Direct memory access (DMA)...

1. reduces the number of interrupts
2. allows for data transfer without the need of interrupts

A page fault could be directly handled by the MMU via DMA. What do you think?

1. why not?
2. not really...

The I/O completion order for device D can differ from how PCBs are sorted in D waiting queue

1. Yes
2. no