

Virtualization Operating Systems – EDA093/DIT401

Vincenzo Gulisano
vincenzo.gulisano@chalmers.se



UNIVERSITY OF
GOTHENBURG

Reading instructions

- 4th Edition: Chapter 7, 7.1-7.7 and 7.9-7.11 (included)
- 5th Edition: Chapter 7, 7.1-7.10 (included)
 - Note: part on Containers only in 5th edition. Slides are enough if you do not have the 5th edition


Agenda

- Introduction
- Requirement for virtualization
- Type 1 and type 2 hypervisors
- Techniques for Efficient Virtualization
- Memory virtualization
- I/O virtualization
- Other benefits of virtualization
- Clouds
- Containers


Agenda

- Introduction
- Requirement for virtualization
- Type 1 and type 2 hypervisors
- Techniques for Efficient Virtualization
- Memory virtualization
- I/O virtualization
- Other benefits of virtualization
- Clouds
- Containers

Without virtualization



With virtualization
(discussed since the 60s...)



Without virtualization

PROS

- If one machine cannot handle the entire load, it works
- Reliability (1 server crash does not imply all services are down)
- Better for security (sandboxing) – if an intruder compromises web server, does not have access to mails
- Run different OSes
- ...

CONS

- Costs money and time

With virtualization

- If the failures are at software/OS level, still reliable
- Run different OSes
- Hypervisor code (which need to be the one you can really trust) is smaller than OS and easier to maintain
- Costs less
- Cheaper to try out applications/services/... (good in software development, to test with different OSes) and provide **virtual appliances**
- Legacy applications
- Easier to migrate services/OSes

- Hardware failure can have catastrophic consequences

As you can imagine, the Cloud starts from this idea and takes it to that of using a server located somewhere else, which further reduces costs...

Agenda

- Introduction
- **Requirement for virtualization**
- Type 1 and type 2 hypervisors
- Techniques for Efficient Virtualization
- Memory virtualization
- I/O virtualization
- Other benefits of virtualization
- Clouds
- Containers

Aside from “acting like a physical machine”, what are actually the requirements?


- Safety (hypervisor should have full control of the virtualized resources)
- Fidelity (program running on bare hardware = program running on a VM)
- Efficiency

Why can't we "just" use an interpreter?

- Wrap each instruction
 - Execute if safe (e.g., INC instruction)
 - Simulate otherwise (e.g., disable interrupts)
- Theoretically it works, but
 - Safety ✓
 - Fidelity ✓
 - Efficiency ✗
- VMMs should execute most of the code directly...

What is required for VMMs to run code directly?

- Sensitive instructions: can be run both in user mode and kernel mode
 - Their behavior might change depending on the case
- Privileged instructions: will result in an interrupt if not in kernel mode



This has been taken into account since 2005

- Both from Intel and AMD
- VT technology:
 - The guest OS runs directly on the hardware until it generates an interrupt
 - Which is delivered to the hypervisor
- The hypervisor sets a hardware bitmap to specify all interrupts it wants to be alerted about (trap-and-emulate approach)
- The technique used before that was binary translation:
 - Simply put: replace all unsafe I/O with a trap




Hardware and
OSes collaborating,
again...


Agenda

- Introduction
- Requirement for virtualization
- **Type 1 and type 2 hypervisors**
- Techniques for Efficient Virtualization
- Memory virtualization
- I/O virtualization
- Other benefits of virtualization
- Clouds
- Containers

Type 1 and Type 2 Hypervisors



(a)




(b)


Agenda

- Introduction
- Requirement for virtualization
- Type 1 and type 2 hypervisors
- **Techniques for Efficient Virtualization**
- Memory virtualization
- I/O virtualization
- Other benefits of virtualization
- Clouds
- Containers

Hypervisor type 1, with VT




Hypervisor type 1, before/without VT



- Before VT, x86 architectures supported 4 protection modes/rings.
- Without virtualization, OS runs in ring 0, user processes in ring 3
- With virtualization, as shown in figure
- OS privileged instructions will trap, but what about sensitive ones?

Hypervisor type 1, before/without VT



- Sensitive instructions are removed by rewriting the code
- Before executing each **basic block** (a straight sequence of instructions that ends with a branch), the hypervisor scans the block and replaces sensitive instructions


The cost of virtualization

- Traps are expensive
- Because of that, sometimes VT + trap-and-emulate performs worst than binary translation
- In fact, some hypervisor might use binary translation even for privileged instructions
- A small example: Suppose you want to disable interrupts for a while
 - VT: set registers in the hardware
 - Binary translation: rewrite code with an extra if statement

Agenda

- Introduction
- Requirement for virtualization
- Type 1 and type 2 hypervisors
- Techniques for Efficient Virtualization
- **Memory virtualization**
- I/O virtualization
- Other benefits of virtualization
- Clouds
- Containers

Why do we need to virtualize memory too?



Shadow page table

- The hypervisor maintains a **shadow page table**, to translate each VM's believed "frame" into the real physical frame
- It is complicated, though:
 - Guest OS page tables and the Hypervisor shadow page table must always be in synch,
 - But the OS can change its page table by simply writing to memory...
 - So, how does the Hypervisor detect a change?
 - Option 1: mark as read-only the VM pages that contain the page table
 - Option 2: do nothing and wait for page faults to happen
 - but this requires extra care of mappings that are removed by the guest OS, because those do not generate page-faults

Page faults

- Guest-induced page faults
 - This are detected by the hypervisor
 - But need to be reinjected to the guest OS
 - Example: OS trying to access a page that has been swapped out of RAM
- Hypervisor-induced page faults
 - Used to keep the shadow page table in synch with the guest OS page table
- When a page faults results in the hypervisor gaining control, we use the term **VM exit**

Reclaiming memory

- If the hypervisor “lies” to the guest OSes pretending it has way more frames it has,
 - and then needs to reclaim some of them from OS A and give them to OS B,
 - how can it do that?
-
- If it simply takes them away, it needs to update the OS page tables, but, where are they? When can that be done safely?...

Ballooning

- A small module is loaded in each VM (a balloon) as a pseudo device driver connected to the hypervisor
- The hypervisor can inflate is requested space to force the guest OS in releasing pages from other processes and give them to this module
 - So that the hypervisor can then use them...

Agenda

- Introduction
- Requirement for virtualization
- Type 1 and type 2 hypervisors
- Techniques for Efficient Virtualization
- Memory virtualization
- **I/O virtualization**
- Other benefits of virtualization
- Clouds
- Containers

I/O virtualization

- The hypervisor probes devices at startup and
- reports them to the guest OSes.
- When the latter try to use them,
- they will interrupt and the hypervisor can intervene
- BUT: each OS cannot really hold the entire device (e.g., a disk)
- Possible solutions
 - Emulate a disk by means of a file managed by the Hypervisor
 - Emulate network by having the Hypervisor actually behaving as a switch

Single Root I/O Virtualization (SR-IOV)

- For performance, a hypervisor-in-the-middle virtualization approach might not always be optimal
- Modern hardware supports virtualization, which allows for devices to be directly controlled by a guest VM (bypassing the hypervisor)
- Devices with SR-IOV usually have
 - Independent memory space, interrupts and DMA streams for each VM
 - Each VM can map in its own memory its corresponding memory areas

Agenda

- Introduction
- Requirement for virtualization
- Type 1 and type 2 hypervisors
- Techniques for Efficient Virtualization
- Memory virtualization
- I/O virtualization
- **Other benefits of virtualization**
- Clouds
- Containers

Other benefits of virtualization - 1

- When a multi-CPU / multi-core machine is used,
- each VM can be assigned a specific subset of cores
- This has an important implication:
 - You could be able to write programs that have a parallelism degree (i.e., number of threads or co-operating processes) that fits the actual number of physical cores the VM you are going to use has access to

Other benefits of virtualization - 2

- Different VMs share memory, and they could potentially be duplicating a lot of (read-only) information
 - Think for instance at having many VMs running the same OS in parallel...
- Deduplication techniques (also referred to as transparent page sharing or content-based page sharing) which are usually used for disks, can be used for memory too!
 - It should not come as a surprise, the physical / logical separation behind VM that is beneficial for running several processes at the same time efficiently can be also beneficial to run several OSes at the same time
 - At the end of the day, the OS runs a collection of processes too

Agenda

- Introduction
- Requirement for virtualization
- Type 1 and type 2 hypervisors
- Techniques for Efficient Virtualization
- Memory virtualization
- I/O virtualization
- Other benefits of virtualization
- **Clouds**
- Containers

What exactly is the cloud?

- Public vs private...
- Access to physical hardware, access to virtual OSeS, access to specific software
- Big machines / small machines
- Desired characteristics
 - On-demand self-service
 - Broad network access (available via network with standard mechanisms)
 - Resource pooling
 - Elasticity
 - Measured service (also in connection to \$)

Types of services

- IAAS (Infrastructure as a Service)
 - Direct access to a specific VM, the user does what (s)he wants
- PASS (Platform as a Service)
 - Specific environment, with given OS, DataBase, WebServer, ...
- SAAS (Software as a Service)
 - Offers specific software

Virtual Machine Migration

- What is the server running a bunch of VMs needs maintenance?
- Intuition: move the VM to a different server...
- Simple approach / but not liked by customers who pay
 - Turn off the VM, move the data to a different server, turn on the VM
 - Long downtime... also, turn off when? Ask the user?
- Live migration
 - Keep copy pages in the back, making the transition from one server to another not noticeable (seamless live migration)
 - As you probably imagine, this is everything but simple...

Checkpointing/Snapshotting

- An advantage of VMs is that they can be paused and copied
- If something bad happens, you can then revert the VM to the previous save
- Incremental snapshotting with mechanisms like copy on write can then allow for lightweight checkpointing
 - Which can be also used in combination with migration

Agenda

- Introduction
- Requirement for virtualization
- Type 1 and type 2 hypervisors
- Techniques for Efficient Virtualization
- Memory virtualization
- I/O virtualization
- Other benefits of virtualization
- Clouds
- **Containers**

Containers (motivation/intuition)

- VMs virtualize hardware
- Containers provide OS-level virtualization
 - Each process / set of processes believes it's running in its own OS
- Note, under the hood:
 - Processes share the same OS and hardware. Processes are isolated, but the performance of a set of processes is always somehow affected by what other processes do
 - All processes / sets of processes need to share OS

Containers – What are the ingredients for a container / a virtualized OS?

- File system
- Process identifiers
- User identifiers
- Network interfaces and IP addresses
- Inter-Process Communication (IPC) endpoints (e.g., pipes)
- Memory limits
- CPU limits
- I/O limits
- ...

Containers – What are the ingredients for a container / a virtualized OS?

- File system
- Process identifiers
- User identifiers
- Network interfaces and IP addresses
- Inter-Process Communication (IPC) endpoints (e.g., pipes)
- Memory limits
- CPU limits
- I/O limits
- ...



Can be managed through namespaces

Of a different nature, they require some active monitoring

Example of namespace: Mount namespace

- Gives a process its own view of the filesystem mount hierarchy
- Processes in different mount namespaces can see different mounted filesystems
- Created or managed using mechanisms such as `clone()`, `unshare()`, and `setns()`
- Container runtimes combine mount namespaces with mounts and mechanisms such as `pivot_root()` or `chroot()` to provide an isolated root filesystem

What about resources that need to be monitored? → cgroups

- One controller per resource type:
 - Resource 1 → cgroups (in a tree) specify “restrictions”
 - Resource 2 → cgroups (in a tree) specify “restrictions”
 - ...
- Each process belongs to exactly one cgroup for each resource of interest, and each resource is managed by a controller
- Example:
 - CPU: c1:50%
 - Memory: m1:2GB
 - Processes A and B should use 50% of CPU: A → c1 in CPU / B → c1 in CPU
 - Process B should not use more than 2GB of memory: B → m1 in Memory

Why are cgroups organized in a tree data structure for a given controller? Why multiple controllers?

- Processes A and B should use 50% of CPU
 - **But what if we want specific quotas for them?**
- Also, assume we have a CPU with 4 cores
 - Pinning A and B to 2 cores → Processes A and B cannot use more than 50% of CPU

Note: there can be different ways of enforcing a restriction, and they might also conflict...

- Cgroup for cpuset controller: use only core 1
 - Affinity of a thread in that cgroup: use only core 2
- The OS will enforce the intersection of the two
- The thread is effectively blocked


cgroups are organized by controller (subsystem)

- cpu → CPU bandwidth or relative CPU weight
- cpuset → which CPU/NUMA part
- memory → memory usage
- io/blkio → disk I/O
- pids → limits the number of processes/tasks
- devices → access to devices

cgroup v1 and cgroup v2


- Please note, the following examples mix v1 and v2 of cgroup
- **cgroup v1**
 - Separate hierarchies may be used for different controllers
 - A process can belong to different cgroups for different controllers
- **cgroup v2**
 - One unified cgroup hierarchy
 - A process belongs to one cgroup in that hierarchy
 - Controllers are enabled within the hierarchy

cgroups are organized by controller (subsystem)



- cpu is defined by means of shares
- Usually:
 - system.slice for system services
 - user.slice for user processes
 - machine.slice for VMs/Containers
- In the example, under full contention, G1 receives approximately 0.25 of the CPU time, G1.A gets $1/11 * 0.25 = 0.02$ of resources

cgroups are organized by controller (subsystem)



- memory is defined in absolute value
- Processes in G2.1 cannot use more than 1 GB
- Processes in G2.2 cannot use more than 1.5 GB
- Processes in G2.1 together with processes in G2.2 cannot use more than 2GB

Each process belongs to one cgroup in each hierarchy

```
cat /proc/4042126/cgroup
```

12:cpuset:/	→ cpuset controller, in the root group (/)
11:perf_event:/	→ performance monitoring events, root group
10:devices:/user.slice	→ limited by device access rules of user.slice
9:memory:/user.slice/user-1009.slice/session-9277.scope	→ memory limits per user/session
8:cpu,cpuacct:/user.slice	→ CPU shares/usage accounting under user.slice
7:blkio:/user.slice	→ block I/O controller, user.slice group
6:rdma:/	→ RDMA resources, unrestricted (root)
5:pids:/user.slice/user-1009.slice/session-9277.scope	→ PID count limits per session
4:hugetlb:/	→ hugepage allocation control, root group
3:net_cls,net_prio:/	→ network classification, root
2:freezer:/	→ process freezing, root
1:name=systemd:/user.slice/user-1009.slice/session-9277.scope	→ systemd bookkeeping
0:./user.slice/user-1009.slice/session-9277.scope	→ cgroup v2 unified hierarchy

You can find information about each controller

```
ls -la /sys/fs/cgroup/blkio/user.slice/
```

```
total 0
```

```
drwxr-xr-x 2 root root 0 Oct 10 14:20 .
```

```
dr-xr-xr-x 5 root root 0 Jul 18 21:47 ..
```

```
--w----- 1 root root 0 Oct 14 06:22 blkio.reset_stats
```

```
-r--r--r-- 1 root root 0 Oct 14 06:22 blkio.throttle.io_service_bytes
```

```
-r--r--r-- 1 root root 0 Oct 14 06:22 blkio.throttle.io_service_bytes_recursive
```

```
-r--r--r-- 1 root root 0 Oct 14 06:22 blkio.throttle.io_serviced
```

```
-r--r--r-- 1 root root 0 Oct 14 06:22 blkio.throttle.io_serviced_recursive
```

```
-rw-r--r-- 1 root root 0 Oct 14 06:22 blkio.throttle.read_bps_device
```

```
-rw-r--r-- 1 root root 0 Oct 14 06:22 blkio.throttle.read_iops_device
```

```
-rw-r--r-- 1 root root 0 Oct 14 06:22 blkio.throttle.write_bps_device
```

```
-rw-r--r-- 1 root root 0 Oct 14 06:22 blkio.throttle.write_iops_device
```

```
-rw-r--r-- 1 root root 0 Oct 14 06:22 cgroup.clone_children
```

```
-rw-r--r-- 1 root root 0 Oct 10 16:53 cgroup.procs
```

```
-rw-r--r-- 1 root root 0 Oct 14 06:22 notify_on_release
```

```
-rw-r--r-- 1 root root 0 Oct 14 06:22 tasks
```

cat /sys/fs/cgroup/blkio/user.slice/blkio.throttle.io_service_bytes
would show the accumulated bytes read/written by this cgroup.

